

Concurrent and Distributed Erlang

Lectures F6 + F7 (Chapters 11..14)

Joe Armstrong

Ericsson AB

W6/F6 - Concurrent Erlang


- Chapters 11,12.
- Concurrency Intro.
- Concurrency Primitives.
- Client Server.
- Processes are cheap.
- Send/Receive.
- Spawn.
- Registered Processes .

Concurrency Intro


- The World is Parallel.
- Concurrency should be in the programming language not the operating system.
- Concurrent is not parallel.

Explain it to a five year old

Concurrent = Two Queues One Coffee Machine



Parallel = Two Queues Two Coffee Machines



© Joe Armstrong 2013

Concurrent vs Parallelism

- Concurrent = *the illusion of parallelism*.
- Parallel = *really parallel*.
- On a quad core there can only really be 4 things happening at the same time.

Note that in reality there is more parallelism than is implied by the number of cores:

- Instruction level parallelism (several assembly instructions can execute at the same time).
- Memory fetch/store parallelism (several memory locations can be transferred to registers at the same time).
- Pipeline parallelism (several things can happen in the pipeline at the same time).
- I/O parallelism (SSD's can have several controllers).

All the above are very difficult to program.

Concurrency Primitives

- spawn, send, receive

```
Pid = spawn(fun() -> ... end)
Pid = spawn(Mod, Func, [Arg1, Arg2, ...])

Pid ! Value

receive
  Pattern1 -> Actions1;
  Pattern2 -> Actions2;
  ...
end
```

- receive is similar to case

```
case Value of
  Pattern1 -> Actions1;
  Pattern2 -> Actions2;
  ...
end
```

Exercise

- Create a registered server called double.
- If you send it an integer it doubles it and sends back the reply.
- It crashes if you send it an atom.
- Make a process that sleeps for a random time and sends a message to the double server and causes it to crash.
- Make a monitor process that detects that the server has crashed. It restarts the server after a random delay.
- Make a client function that sends a request to the server and times out if the request is not satisfied. We can assume the server has crashed. The client should wait a second and then try again.
- Abort the client if it has tried more than ten times.

Client - Server 1

```
start1() ->  
  spawn(f6, loop1, []).
```

```
loop1() ->  
  receive  
    {square, X} ->  
      print(X*X),  
      loop1()  
  end.
```

```
Pid ! {square, 10}
```

- Run with `f6:start1()`.
- **Pid** always means **Process Identifier**.
- How do we get the result back?

Client - Server 2

- Getting the result back?

```
start2() ->
  spawn(f6, loop2, [])

loop2() ->
  receive
    {From, {square, X}} ->
      From ! X*X,
      loop2()
  end.

Pid ! {self(), {square, 10}},
receive
  Result ->
    ...
end
```

- How do we know the result we got back was from the server and not from some other process that just happened to send us a message?

Client - Server 3

- Pattern match the reply message to check that the reply comes from the correct process.

```
start3() ->
  spawn(f6, loop3, []).

loop3() ->
  receive
    {From, {square, X}} ->
      From ! {self(), X*X},
      loop3()
  end.

area_square(Pid, X) ->
  Pid ! {self(), {square, X}},
  receive
    {Pid, Area} -> Area
  end.
```

- Say something about selective receive ...

Selective Receive

```
receive
  Pattern1 ->
    Actions1;
  Pattern2 ->
    Actions2
end
```

This suspends until a message matching `Pattern1` or `Pattern2` is received. **All other messages are queued.**

Client - Server 4 (abstract the RPC)

```
%% old

area_square(Pid, X) ->
  Pid ! {self(), {square, X}},
  receive
    {Pid, Area} -> Area
  end.

%% refactored

area_square(Pid, X) ->
  rpc(Pid, {square, X}).

rpc(Pid, Query) ->
  Pid ! {self(), Query},
  receive
    {Pid, Reply} ->
      Reply
  end.
```

Client - Server 5 (Tagged replies)

```
Pid = spawn(fun() -> loop() end)

loop() ->
  receive
    {From, Tag, {square, X}} ->
      Result = X*X,
      From ! {Tag, Result},
      loop()
  end.

rpc(Pid, Query) ->
  Tag = erlang:make_ref(),
  Pid ! {self(), Tag, Query},
  receive
    {Tag, Result} ->
      Result
  end.
```

Timeouts

```
receive
  Pattern1 ->
    Actions1;
  Pattern2 ->
    Actions2;
  ...
after Time ->
  Actions
end.
```

Client - Server 6

- We detect that the server has not replied with a timeout...

```
rpc(Pid, Query) ->
    Tag = erlang:make_ref(),
    Pid ! {self(), Tag, Query},
    receive
        {Tag, Result} ->
            {ok, Result}
        after 1000 ->
            {error, timeout}
    end.
```

Client - Server 7

- Umm

```
rpc(Pid, Query) ->  
    Tag = erlang:make_ref(),  
    Pid ! {self(), Tag, Query},  
    receive  
        {Tag, Result} ->  
            Result  
    after TIME ->  
        DO SOMETHING  
    end.
```

- What is DO SOMETHING?
- What is TIME?
- Idempotence.
- Getting DO SOMETHING and TIME right is incredibly difficult.

Why is this difficult?

- We send a message to a server.
- We do not get a reply

So:

- Either the server has crashed, or,
- The communication channel is broken.

Recovering from this is very difficult (in many cases it is impossible).

Exercise (reminder)

DO SOMETHING means:

- Try again N times with a random delay and then give up.
- Write some code to randomly crash the server.
- Write some code to restart the server if it crashes.

What really happens - the mailbox

- Each process has a mailbox.
- Send causes a message to be added to the mailbox.
- When a process message is added to a mailbox the process is scheduled for execution.
- When the process next executes it checks if the new mails match any of the receive patterns.
- If the message does not match the process suspends.

The scheduler

- Processes run for 1000 reductions and are then suspended. They stay in the run queue.
- Processes waiting for a message are removed from the run-queue.
- When a message is added to the mailbox we add it to the run queue (if it is not in the run queue).
- There is one sheduler per core *not really true – can be two or more*.
- Processes can be moved between schedulers.



The scheduler

Client Server patterns

```
Pid = spawn(fun() -> loop(State) end)
```

```
loop(State) ->
```

```
  receive
```

```
    {From, Pattern1} ->
```

```
      State1 = ...
```

```
      Result = ...
```

```
      From ! {self(), Result},
```

```
      loop(State1);
```

```
    {From, Pattern2} ->
```

```
      ...
```

```
  end.
```

```
funcl(Pid, Args) -> rpc(Pid, Args).
```

```
rpc(Pid, Args) ->
```

```
  Pid ! {self(), Args},
```

```
  receive
```

```
    {Pid, Ret} -> Ret
```

```
  end.
```

A Stateful counter

```
Pid = spawn(fun() -> counter(0) end)
```

```
counter(N) ->
  receive
    {From, {add,K}} ->
      From ! {self(), ok},
      counter(N+K)
  end.
```

```
add(K) -> rpc(Pid, {add,K}).
```

```
rpc(Pid, Msg) ->
  Pid ! {self(), Msg},
  receive
    {Pid, Reply} ->
      Reply
  end.
```

Extend the server

```
receive
  ...

  {From, reset} -
    counter(0)
  ...
  {From, read} ->
    From ! {self(), N},
    counter(N);
  ...
end.

reset(K) -> rpc(Pid, reset).
...
```

- Add extra patterns in the server.
- Add API routines.

Send functions in the messages

```
Pid = spawn(fun() -> loop(State) end)

loop(State) ->
  receive
    {From, F} ->
      {Reply, NewState} = F(State),
      From ! {self(), Reply},
      loop(NewState)
  end.

add(K) -> rpc(Pid,
  fun(State) ->
    {ack, K+State}
  end).
```

Send the server in a message

```
start() ->
  spawn(fun() -> wait() end)

wait() ->
  receive
    {become, F} ->
      F()
  end.

Pid = start(),
...
Pid ! {become, fun() -> loop/1}.

loop(State) ->
  receive
    ...
  end
```

processes are cheap

```
-module(f6).  
-compile(export_all).  
  
time(N) ->  
    {Time, _} = timer:tc(f6, time_test, [N]),  
    Tsec = Time / 1000000,  
    {spawned, trunc(N / Tsec), 'processes/sec'}.  
  
time_test(0) ->  
    true;  
time_test(N) ->  
    spawn(fun() -> true end),  
    time_test(N-1).
```

- show this.

erl -smp disable

```
erl -smp disable
Eshell V5.10.1 (abort with ^G)
1> f6:time(1000000).
{spawned,1027305,'processes/sec'}
2> f6:time(1000000).
{spawned,1212416,'processes/sec'}
```

- 1.2 Million processes/sec.

Registered Processes

- `Pid ! Message` sends a message to the mailbox of the process `Pid`.
- How do we know `Pid`?
- Only the parent knows `Pid`

```
start() ->  
    Pid = spawn(...),  
    Pid ! Message,  
    ...
```

Registered Processes

```
start() ->  
  Pid = spawn(...),  
  register(counter, Pid),  
  ...
```

- Now any process can send a message to the process

```
> counter ! {add, 12}
```

Tail recursion

```
start() -> spawn(Mod, loop, [Arg1, ...]).  
  
loop(Arg1, ...) ->  
  receive  
    Pattern1 ->  
      ...  
      loop(1);  
    Pattern2 ->  
      ...  
  end
```

- The last thing you do is call yourself.

Non Tail recursion

```
start() -> spawn(Mod, loop, [Arg1, ...]).

loop(Arg1, ...) ->
  receive
    Pattern1 ->
      ...
      loop(1, ..), %% NO NO NO NO
      ... <-- Don't call stuff after ... ; the call to loop
    Pattern2 ->
      ...
  end
```


Tail recursion (again)

- Co-routines.
- Continuation passing style.

```
state1(...) ->
  receive
    Pattern1 ->
      ...
      state2(1);
    Pattern2 ->
      ...
  end.

state2(...) ->
  receive
    Pattern1 ->
      ...
      state3(...);
    ...
  end
```

- If something never returns, it must be the last thing you call.

What does Mod:Func really mean?

- What's the difference between `loop` and `loop1`?

```
-module(foo).  
  
loop(State) ->  
  receive  
    Pattern1 ->  
      ...  
      loop(NewState)  
  end.  
  
loop1(State) ->  
  receive  
    Pattern1 ->  
      ...  
      foo:loop1(NewState)  
  end.
```

- **Mod:Func calls the latest version.**

Spawn MFA - or fun

```
start1() ->  
  spawn(Mod, Func, [Arg1, Arg2, ..., ArgN])  
  
start2() ->  
  spawn(fun() -> ... end)
```

W6/F7 - Error and Distributed Programming

- Chapters 13,14.
- Links.
- Monitors.
- The error model.
- Firewall.
- Generic Allocator.
- Why Distributed Programming?
- Erlang distribution.
- Explicit sockets and protocols.
- Distribution BIFS.
- Nano-twitter.
- Security.

Links

- When an unlinked process dies nobody will know about it.
- When a linked process dies the processes in the link-set of the process will be sent an error signal.
- A normal process receiving a non-normal error signal will die.
- A system process will convert the error signal to an error message and can receive it like any other message.
- A processes that terminates normally sends a normal error signal to its link set.

signal: {'EXIT', Pid, Why}

Why kan vara: normal

kan också vara: kill

Error handling BIFs

- `link(Pid) unlink(Pid)`
- `process_flag(trap_exit, true)`
- `spawn_link(Fun) spawn_link(M, F, A)`
- `exit(Pid, Why) exit(Pid, kill)`

Trapping exits

```
start() ->
  spawn_link(M, F, A) end.

start() ->
  process_flag(trap_exit, true),
  loop().

loop() ->
  receive
    {'EXIT', Pid, Why} ->
      ...
  end.
```

A process that restarts a failed process

```
start_and_watch() ->
    Pid = spawn(f6, loop1, []),
    spawn(f6, watch, [Pid]),
    Pid.

watch(Pid) ->
    link(Pid),
    process_flag(trap_exit, true),
    print({watching, Pid}),
    receive
        {'EXIT', Pid, Why} ->
            print({ohDear, Pid, crashed, because, Why})
    end.
```


Running the watcher

```
Pid2 = f6:start_and_watch().
{watching,<0.62.0>}
<0.62.0>
9> Pid2 ! {square,abc}.
{ohDear,<0.62.0>,crashed,because,
    {badarith,[{f6,loop1,0,[{file,"f6.erl"},{line,21}]}}]}

=ERROR REPORT==== 31-Jan-2014::10:26:29 ===
Error in process <0.62.0> with exit value:
    {badarith,[{f6,loop1,0,[{file,"f6.erl"},{line,21}]}}]}

{square,abc}
```

Spawn and Link race conditions

```
start_and_watch() ->
    Pid = spawn(f6, loop1, []),
    spawn(f6, watch, [Pid]),
    Pid.

watch(Pid) ->
    link(Pid),
    process_flag(trap_exit, true),
    ...
```

- If you are very quick, loop1 might crash *before* the watcher executes the `link` statement.
- That's why we have `spawn_link`.

Monitors

- One directional links.
- A monitor is like two links “back to back”.

monitors: ----->

links: <----->

The Worker/Manager model

- Workers do work. They crash if they cannot do what they are supposed to do.
- Managers detect the failure of workers and restart them.
- `start_and_watch` was a very simple Worker/Manager program.
- OTP has supervisors which generalises this idea.
- Akka (Java) is a clone of this idea.

Generic Allocator

- Server links to clients and deallocates resources if the clients crash

```
loop(State) ->
  receive
    {From, Tag, {allocate, X}} ->
      {Resource, State1} = allocate(From, X, State),
      link(From),
      From ! {Tag, Resource},
      loop(State1);
    ...

    {'EXIT', Pid, _} ->
      State1 = dealloc_resources_owned_by(Pid, State),
      loop(State1)
  end.
```

Distributed Programming?

- Fun.
- World is distributed.

Why? ITS FUN =]

Erlang distribution

- Needs only one new idea.

extra

- `spawn(Node, Mod, Func, Args).`
- Remote Pids work just like local Pids.
- Can test on one machine, deploy on many.

Nano twitter

```
start() ->
    register(twit, spawn(nano_twitter, watcher, [])).

watcher() ->
    receive
        Any ->
            print({tweet, Any}),
            watcher()
        after 5000 ->
            print(yawn),
            watcher()
    end.

connect() ->
    pong = net_adm:ping('twit@joe').

tweet(Msg) ->
    rpc:cast('twit@joe', erlang, send, [twit, Msg]).
```


Running Nano Twitter

In one terminal:

```
erl -noshell -sname twit -s nano_twitter start  
yawn  
...
```

Someplace else:

```
$ erl -sname one  
(one@joe)1> nano_twitter:connect().  
pong  
(one@joe)2> nano_twitter:tweet('hi joe').  
true
```

Back where you started:

```
yawn  
{tweet, 'hi joe' }
```

Fun Exercise (1)

- No credits.
- Extend `nano_twitter`
- Work in pairs.
- Client on one machine.
- Same LAN.
- Server on a different machine

Fun Exercise (2)

- No credits.
- Server on a different machine with a different OS.
- On a WAN.

Fun Exercise (3)

- No credits.
- Get the entire class running.

Fun Exercise (4)

- Credits.
- \$\$\$\$
- Quit school.
- Form a company.
- Connect the world together.

RTFM



http://www.erlang.org/doc/reference_manual/distributed

- Designed for clusters in the same LAN not wide scale distribution.
- Cookie security.
- Great for tightly connected clusters in the same administrative domain.
- Code distribution problem not solved (assumes same backend to fetch the code from) - originally all nodes read from a NFS file system.

Next Week

- Doing it with sockets :-)

Have fun

